

Homework 3 - Chapter 4

Sherwyn Braganza, Mangesh Sakordekar

This is the initial copy of code from Chapter 4

```
In [1]: from tensorflow.keras.datasets import imdb
        (train_data, train_labels), (test_data, test_labels) = imdb.load_data(
            num_words=10000)
```

```
In [2]: max([max(sequence) for sequence in train_data])
```

```
Out[2]: 9999
```

```
In [3]: word_index = imdb.get_word_index()
        reverse_word_index = dict(
            [(value, key) for (key, value) in word_index.items()])
        decoded_review = " ".join(
            [reverse_word_index.get(i - 3, "?") for i in train_data[2]])
```

```
In [4]: import numpy as np
        def vectorize_sequences(sequences, dimension=10000):
            results = np.zeros((len(sequences), dimension))
            for i, sequence in enumerate(sequences):
                for j in sequence:
                    results[i, j] = 1.
            return results

        # x corresponds to the data
        x_train = vectorize_sequences(train_data)
        x_test = vectorize_sequences(test_data)
        print(x_train.shape)
```

```
(25000, 10000)
```

```
In [5]: # y corresponds to the labels
        y_train = np.asarray(train_labels).astype("float32")
        y_test = np.asarray(test_labels).astype("float32")
        print(y_train.shape)
```

```
(25000,)
```

1) Change the model to use one hidden (aka "representation") layer, and see how this affects validation and test accuracy. Show this by plotting them (on the same plot, like in the book). Repeat this for three hidden layers. Comment on any differences you find, as well as what you'd expect.

Part 1 - One Layer Model

```
In [6]: from tensorflow import keras
        from tensorflow.keras import layers

        model = keras.Sequential([
            layers.Dense(16, activation="relu"),
            layers.Dense(1, activation="sigmoid")
        ])

        model.compile(optimizer="rmsprop",
                       loss="binary_crossentropy",
                       metrics=["accuracy"])

        x_val = x_train[:10000]
        partial_x_train = x_train[10000:]
        y_val = y_train[:10000]
        partial_y_train = y_train[10000:]

        history = model.fit(partial_x_train,
                            partial_y_train,
                            epochs=20,
                            batch_size=512,
                            validation_data=(x_val, y_val))

        history_dict_onelayer = history.history
        history_dict_onelayer.keys()

Epoch 1/20
30/30 [=====] - 6s 153ms/step - loss: 0.5228 - acc
uracy: 0.7941 - val_loss: 0.4145 - val_accuracy: 0.8629
Epoch 2/20
30/30 [=====] - 1s 38ms/step - loss: 0.3386 - accu
racy: 0.8995 - val_loss: 0.3321 - val_accuracy: 0.8803
Epoch 3/20
30/30 [=====] - 1s 26ms/step - loss: 0.2608 - accu
racy: 0.9191 - val_loss: 0.2967 - val_accuracy: 0.8883
Epoch 4/20
30/30 [=====] - 1s 24ms/step - loss: 0.2150 - accu
racy: 0.9321 - val_loss: 0.2834 - val_accuracy: 0.8881
Epoch 5/20
30/30 [=====] - 1s 27ms/step - loss: 0.1811 - accu
```

```

racy: 0.9434 - val_loss: 0.2755 - val_accuracy: 0.8901
Epoch 6/20
30/30 [=====] - 1s 29ms/step - loss: 0.1566 - accu
racy: 0.9530 - val_loss: 0.2781 - val_accuracy: 0.8877
Epoch 7/20
30/30 [=====] - 1s 32ms/step - loss: 0.1370 - accu
racy: 0.9596 - val_loss: 0.2851 - val_accuracy: 0.8840
Epoch 8/20
30/30 [=====] - 1s 35ms/step - loss: 0.1206 - accu
racy: 0.9651 - val_loss: 0.2835 - val_accuracy: 0.8866
Epoch 9/20
30/30 [=====] - 1s 30ms/step - loss: 0.1068 - accu
racy: 0.9703 - val_loss: 0.2921 - val_accuracy: 0.8850
Epoch 10/20
30/30 [=====] - 1s 28ms/step - loss: 0.0952 - accu
racy: 0.9741 - val_loss: 0.3051 - val_accuracy: 0.8830
Epoch 11/20
30/30 [=====] - 1s 27ms/step - loss: 0.0841 - accu
racy: 0.9795 - val_loss: 0.3170 - val_accuracy: 0.8830
Epoch 12/20
30/30 [=====] - 1s 29ms/step - loss: 0.0747 - accu
racy: 0.9827 - val_loss: 0.3330 - val_accuracy: 0.8785
Epoch 13/20
30/30 [=====] - 1s 29ms/step - loss: 0.0661 - accu
racy: 0.9858 - val_loss: 0.3623 - val_accuracy: 0.8767
Epoch 14/20
30/30 [=====] - 1s 25ms/step - loss: 0.0589 - accu
racy: 0.9875 - val_loss: 0.3551 - val_accuracy: 0.8795
Epoch 15/20
30/30 [=====] - 1s 24ms/step - loss: 0.0512 - accu
racy: 0.9895 - val_loss: 0.3695 - val_accuracy: 0.8790
Epoch 16/20
30/30 [=====] - 1s 23ms/step - loss: 0.0454 - accu
racy: 0.9915 - val_loss: 0.3896 - val_accuracy: 0.8724
Epoch 17/20
30/30 [=====] - 1s 24ms/step - loss: 0.0403 - accu
racy: 0.9927 - val_loss: 0.4076 - val_accuracy: 0.8760
Epoch 18/20
30/30 [=====] - 1s 25ms/step - loss: 0.0350 - accu
racy: 0.9945 - val_loss: 0.4230 - val_accuracy: 0.8752
Epoch 19/20
30/30 [=====] - 1s 24ms/step - loss: 0.0314 - accu
racy: 0.9948 - val_loss: 0.4401 - val_accuracy: 0.8710
Epoch 20/20
30/30 [=====] - 1s 23ms/step - loss: 0.0276 - accu
racy: 0.9957 - val_loss: 0.4584 - val_accuracy: 0.8719

```

```
Out[6]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

Results of Model with One Layer

```
In [7]: import matplotlib.pyplot as plt
```

```
# Training Loss vs Validation Loss
```

```
loss_values = history_dict_onelayer["loss"]
```

```
val_loss_values = history_dict_onelayer["val_loss"]
```

```
epochs = range(1, len(loss_values) + 1)
```

```
plt.plot(epochs, loss_values, label="Training loss", color = 'red', linestyle='--')
```

```
plt.plot(epochs, val_loss_values, label="Validation loss", color = 'green', linestyle='--')
```

```
plt.title("One Hidden Layer Training and validation loss")
```

```
plt.xlabel("Epochs")
```

```
plt.ylabel("Loss")
```

```
plt.legend()
```

```
plt.show()
```

```
plt.clf()
```

```
# Training Accuracy vs. Validation Accuracy
```

```
acc = history_dict_onelayer["accuracy"]
```

```
val_acc = history_dict_onelayer["val_accuracy"]
```

```
plt.plot(epochs, acc, label="Training acc", color='red', linestyle='--', marker='^')
```

```
plt.plot(epochs, val_acc, label="Validation acc", color='green', linestyle='--', marker='^')
```


```
plt.title("One Hidden Layer Training and validation accuracy")
```

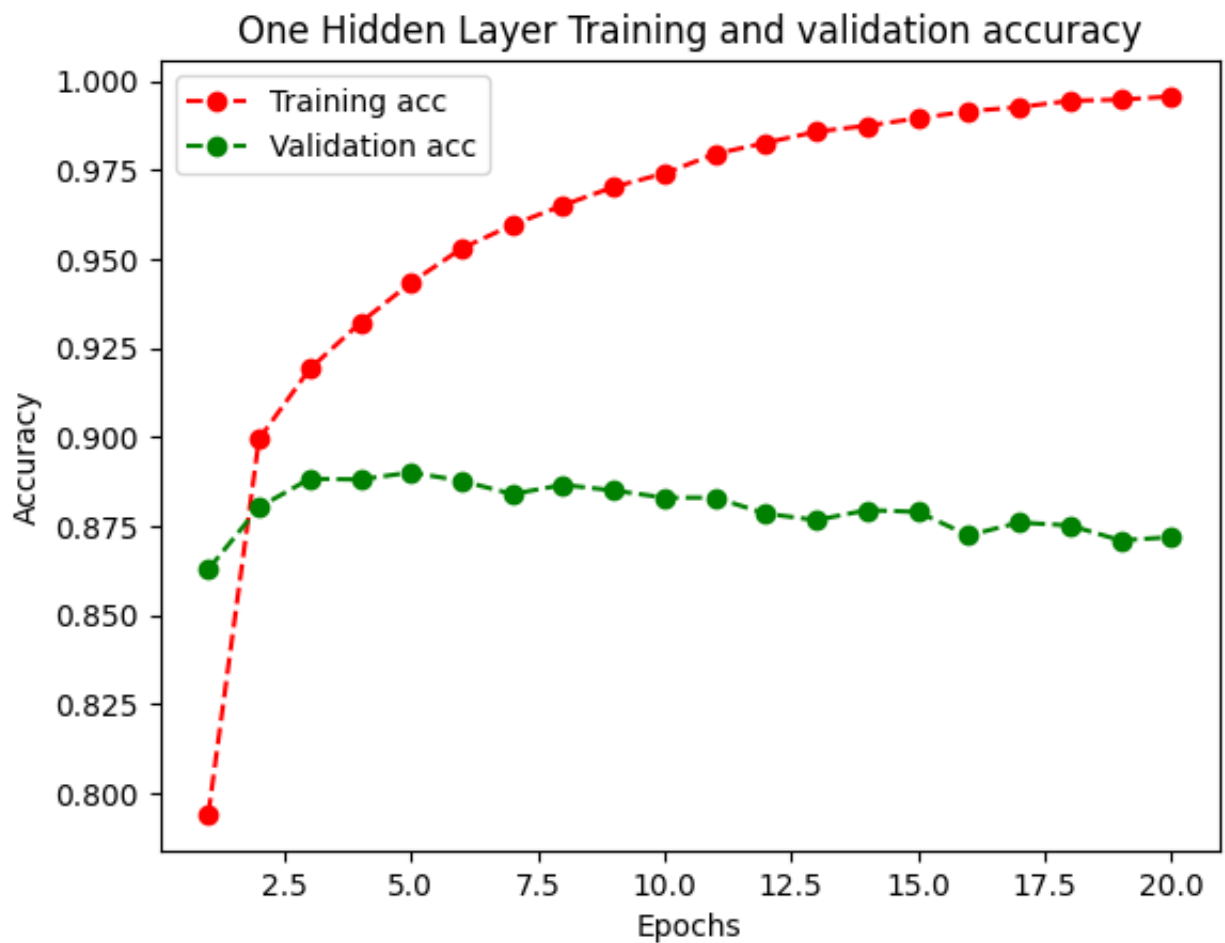
```
plt.xlabel("Epochs")
```

```
plt.ylabel("Accuracy")
```

```
plt.legend()
```

```
plt.show()
```





Part 2 - 3 Layer Model

```
In [8]: model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])

history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))

history_dict_threelayers = history.history
history_dict_threelayers.keys()
```

Epoch 1/20
 30/30 [=====] - 4s 58ms/step - loss: 0.5367 - accuracy: 0.7705 - val_loss: 0.3985 - val_accuracy: 0.8614
 Epoch 2/20

30/30 [=====] - 1s 30ms/step - loss: 0.3120 - accuracy: 0.9000 - val_loss: 0.3268 - val_accuracy: 0.8700
Epoch 3/20
30/30 [=====] - 1s 34ms/step - loss: 0.2210 - accuracy: 0.9277 - val_loss: 0.2784 - val_accuracy: 0.8898
Epoch 4/20
30/30 [=====] - 1s 28ms/step - loss: 0.1713 - accuracy: 0.9431 - val_loss: 0.2884 - val_accuracy: 0.8852
Epoch 5/20
30/30 [=====] - 1s 24ms/step - loss: 0.1357 - accuracy: 0.9568 - val_loss: 0.3007 - val_accuracy: 0.8848
Epoch 6/20
30/30 [=====] - 1s 23ms/step - loss: 0.1105 - accuracy: 0.9657 - val_loss: 0.3252 - val_accuracy: 0.8768
Epoch 7/20
30/30 [=====] - 1s 25ms/step - loss: 0.0928 - accuracy: 0.9703 - val_loss: 0.4247 - val_accuracy: 0.8636
Epoch 8/20
30/30 [=====] - 1s 32ms/step - loss: 0.0770 - accuracy: 0.9778 - val_loss: 0.3520 - val_accuracy: 0.8809
Epoch 9/20
30/30 [=====] - 2s 57ms/step - loss: 0.0565 - accuracy: 0.9849 - val_loss: 0.3877 - val_accuracy: 0.8735
Epoch 10/20
30/30 [=====] - 1s 39ms/step - loss: 0.0457 - accuracy: 0.9884 - val_loss: 0.4302 - val_accuracy: 0.8756
Epoch 11/20
30/30 [=====] - 1s 39ms/step - loss: 0.0352 - accuracy: 0.9916 - val_loss: 0.4565 - val_accuracy: 0.8707
Epoch 12/20
30/30 [=====] - 1s 35ms/step - loss: 0.0313 - accuracy: 0.9914 - val_loss: 0.4842 - val_accuracy: 0.8743
Epoch 13/20
30/30 [=====] - 1s 31ms/step - loss: 0.0228 - accuracy: 0.9947 - val_loss: 0.5177 - val_accuracy: 0.8737
Epoch 14/20
30/30 [=====] - 1s 32ms/step - loss: 0.0180 - accuracy: 0.9955 - val_loss: 0.5566 - val_accuracy: 0.8712
Epoch 15/20
30/30 [=====] - 1s 30ms/step - loss: 0.0083 - accuracy: 0.9991 - val_loss: 0.6149 - val_accuracy: 0.8609
Epoch 16/20
30/30 [=====] - 1s 25ms/step - loss: 0.0168 - accuracy: 0.9956 - val_loss: 0.6381 - val_accuracy: 0.8695
Epoch 17/20
30/30 [=====] - 1s 25ms/step - loss: 0.0041 - accuracy: 0.9997 - val_loss: 0.6736 - val_accuracy: 0.8670
Epoch 18/20
30/30 [=====] - 1s 24ms/step - loss: 0.0136 - accuracy: 0.9963 - val_loss: 0.7203 - val_accuracy: 0.8667
Epoch 19/20
30/30 [=====] - 1s 26ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.7403 - val_accuracy: 0.8665

Epoch 20/20

30/30 [=====] - 1s 25ms/step - loss: 0.0048 - accuracy: 0.9985 - val_loss: 0.8746 - val_accuracy: 0.8615

Out[8]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

In [9]: **import** matplotlib.pyplot **as** plt

Training Loss vs Validation Loss

loss_values = history_dict_threelayers["loss"]

val_loss_values = history_dict_threelayers["val_loss"]

epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, loss_values, label="Training loss", color = 'red', linestyle='--')

plt.plot(epochs, val_loss_values, label="Validation loss", color = 'green', linestyle='--')

plt.title("One Hidden Layer Training and validation loss")

plt.xlabel("Epochs")

plt.ylabel("Loss")

plt.legend()

plt.show()

plt.clf()

Training Accuracy vs. Validation Accuracy

acc = history_dict_threelayers["accuracy"]

val_acc = history_dict_threelayers["val_accuracy"]

plt.plot(epochs, acc, label="Training acc", color='red', linestyle='--', marker='o')

plt.plot(epochs, val_acc, label="Validation acc", color='green', linestyle='--', marker='o')

plt.title("One Hidden Layer Training and validation accuracy")


plt.xlabel("Epochs")


plt.ylabel("Accuracy")

plt.legend()

plt.show()

One Hidden Layer Training and validation loss





One Layer vs 3 Layer

```

In [10]: loss_values = [
            history_dict_onelayer['loss'],
            history_dict_onelayer['val_loss'],
            history_dict_threelayers['loss'],
            history_dict_threelayers['val_loss']
        ]

acc_values = [
            history_dict_onelayer['accuracy'],
            history_dict_onelayer['val_accuracy'],
            history_dict_threelayers['accuracy'],
            history_dict_threelayers['val_accuracy']
        ]

loss_labels = ['One Layer ', 'Three Layer ', 'Training Loss', 'Validation Loss']
accuracy_labels = ['One Layer ', 'Three Layer ', 'Training Accuracy', 'Validation Accuracy']
colors = ['red', 'green', 'blue', 'magenta']

for idx, x in enumerate(loss_values):
    plt.plot(epochs, x, label=loss_labels[int(idx/2)] + loss_labels[idx%2],
             color=colors[idx], linestyle='--', marker = '^')

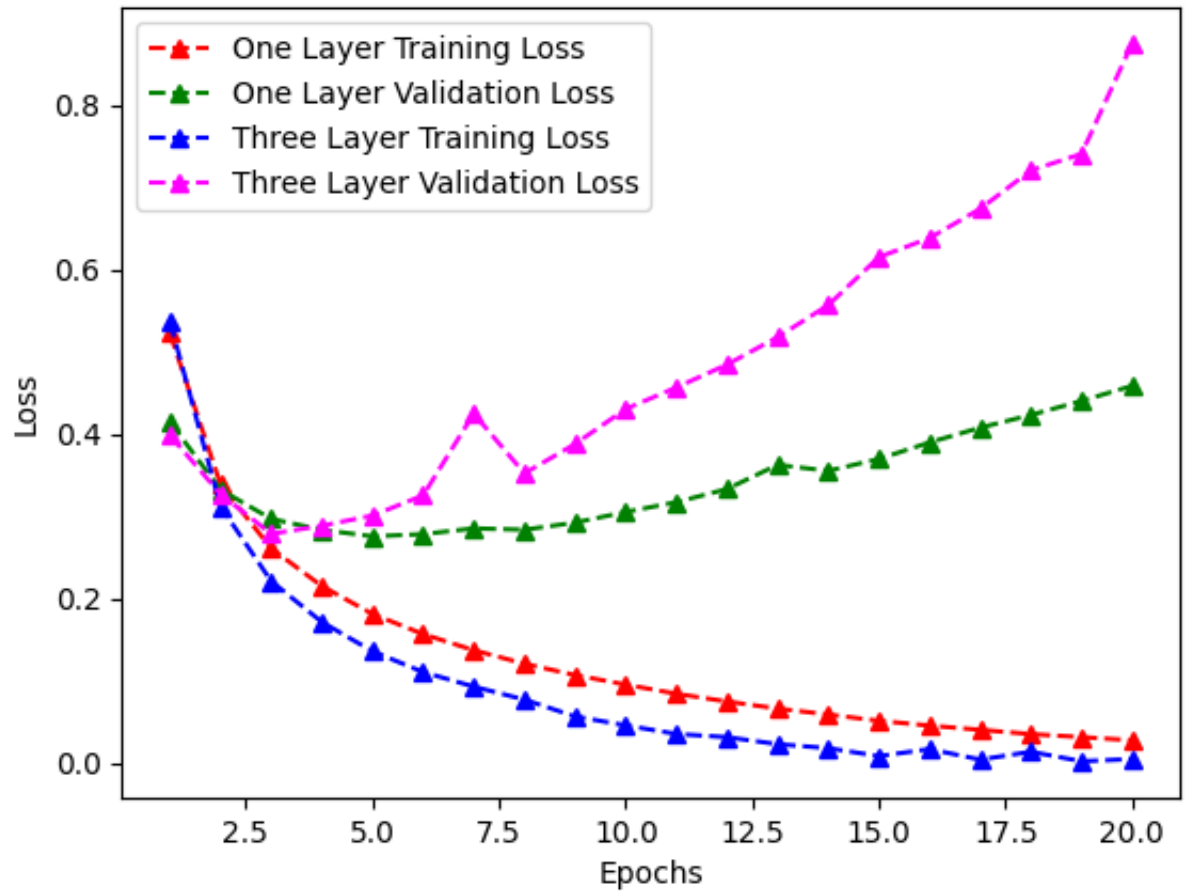
plt.title("One Layer vs. Three Layer Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
plt.clf()

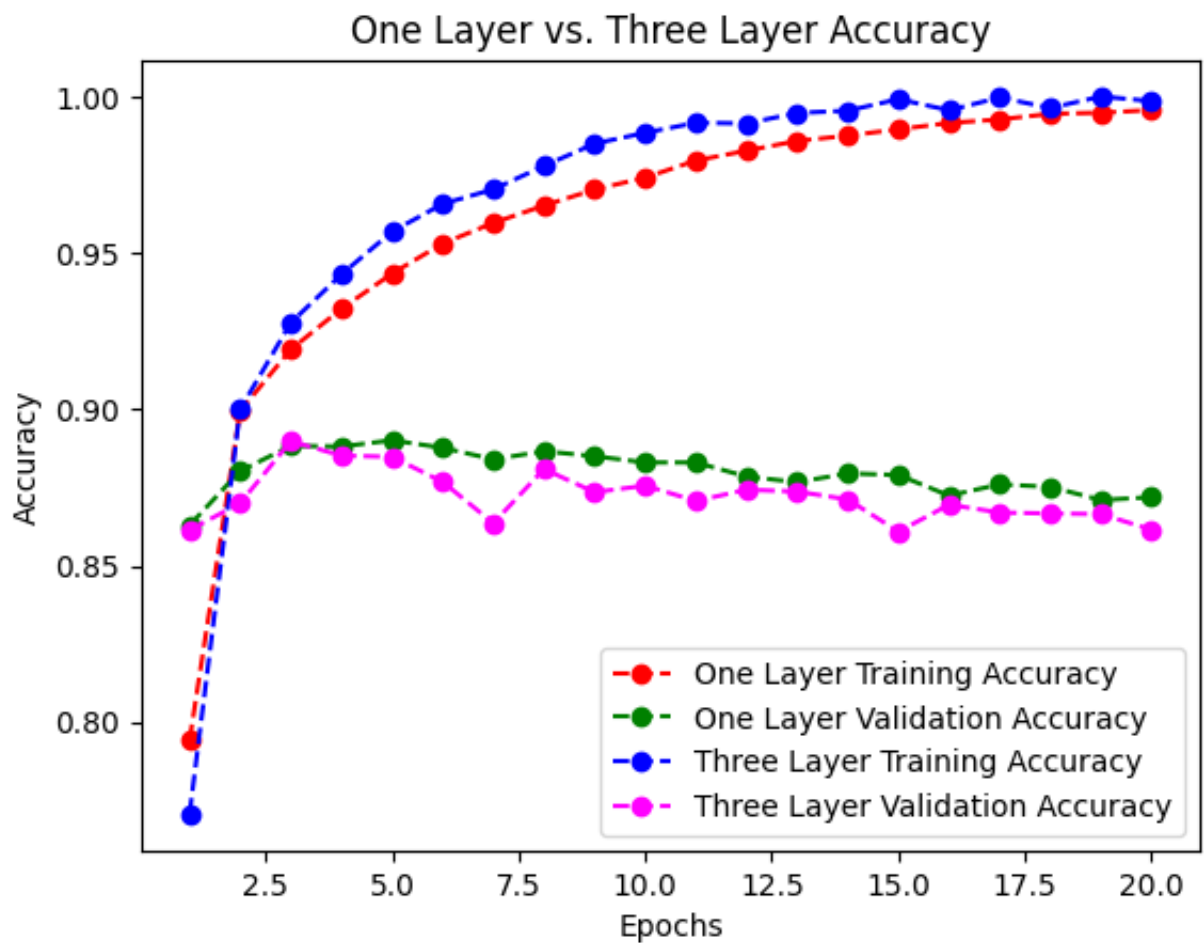
for idx, x in enumerate(acc_values):
    plt.plot(epochs, x, label=accuracy_labels[int(idx/2)] + accuracy_labels[idx%2],
             color=colors[idx], linestyle='--', marker = 'o')

plt.title("One Layer vs. Three Layer Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
plt.clf()

```

One Layer vs. Three Layer Loss





<Figure size 640x480 with 0 Axes>

The best validation accuracy achieved was 88.98

by the 3 layer model. The best validation accuracy by the one layer model was 88.99

On comparing the loss and accuracies of the one layer and three layer side by side, we can see that the validation accuracy changes for the worse after the 4th epoch. Conversely, the accuracy and loss corresponding to the training set steadily increases past the 4th epoch - this implies that the model is overfitting itself to the training data - and from this we can infer that the best model is achieved soon after the 4th epoch of training.

Another key feature to note is that both models have relatively similar loss and accuracy on the Validation Set uptill the 4th epoch, but the Three Layer model exhibits a higher training accuracy and lower loss than the one layer model; thereby telling us that the Three Layer Model performs slightly better.

2) Try changing the number of nodes in the hidden layers to 100. See how this affects validation and test accuracy. Show this by plotting them (on the same plot, like in the book).**

Part 1 - 1 Layer, 100 Nodes

In [11]: *# Model Generation*

```
model = keras.Sequential([
    layers.Dense(100, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])

history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))

history_dict_onelayer = history.history
history_dict_onelayer.keys()
```

Epoch 1/20

30/30 [=====] - 4s 100ms/step - loss: 0.4662 - accuracy: 0.7889 - val_loss: 0.3418 - val_accuracy: 0.8647

Epoch 2/20

30/30 [=====] - 2s 78ms/step - loss: 0.2543 - accuracy: 0.9090 - val_loss: 0.2804 - val_accuracy: 0.8902

Epoch 3/20

30/30 [=====] - 2s 59ms/step - loss: 0.1864 - accuracy: 0.9353 - val_loss: 0.2940 - val_accuracy: 0.8821

Epoch 4/20

30/30 [=====] - 2s 52ms/step - loss: 0.1518 - accuracy: 0.9482 - val_loss: 0.3046 - val_accuracy: 0.8801

Epoch 5/20

30/30 [=====] - 2s 55ms/step - loss: 0.1234 - accuracy: 0.9590 - val_loss: 0.2998 - val_accuracy: 0.8843

Epoch 6/20

30/30 [=====] - 2s 54ms/step - loss: 0.0995 - accuracy: 0.9683 - val_loss: 0.3891 - val_accuracy: 0.8630

Epoch 7/20

30/30 [=====] - 2s 54ms/step - loss: 0.0797 - accuracy: 0.9785 - val_loss: 0.3419 - val_accuracy: 0.8800

Epoch 8/20

30/30 [=====] - 2s 57ms/step - loss: 0.0668 - accuracy: 0.9819 - val_loss: 0.3758 - val_accuracy: 0.8765

```


Epoch 9/20
30/30 [=====] - 1s 50ms/step - loss: 0.0540 - accu
racy: 0.9846 - val_loss: 0.4650 - val_accuracy: 0.8559
Epoch 10/20
30/30 [=====] - 2s 55ms/step - loss: 0.0452 - accu
racy: 0.9886 - val_loss: 0.4081 - val_accuracy: 0.8747
Epoch 11/20
30/30 [=====] - 2s 53ms/step - loss: 0.0339 - accu
racy: 0.9929 - val_loss: 0.5188 - val_accuracy: 0.8595
Epoch 12/20
30/30 [=====] - 2s 55ms/step - loss: 0.0350 - accu
racy: 0.9920 - val_loss: 0.4684 - val_accuracy: 0.8712
Epoch 13/20
30/30 [=====] - 2s 55ms/step - loss: 0.0195 - accu
racy: 0.9976 - val_loss: 0.5065 - val_accuracy: 0.8664
Epoch 14/20
30/30 [=====] - 2s 53ms/step - loss: 0.0180 - accu
racy: 0.9975 - val_loss: 0.6879 - val_accuracy: 0.8474
Epoch 15/20
30/30 [=====] - 2s 55ms/step - loss: 0.0148 - accu
racy: 0.9981 - val_loss: 0.5580 - val_accuracy: 0.8692
Epoch 16/20
30/30 [=====] - 1s 50ms/step - loss: 0.0133 - accu
racy: 0.9985 - val_loss: 0.5927 - val_accuracy: 0.8683
Epoch 17/20
30/30 [=====] - 2s 55ms/step - loss: 0.0127 - accu
racy: 0.9977 - val_loss: 0.6254 - val_accuracy: 0.8660
Epoch 18/20
30/30 [=====] - 2s 56ms/step - loss: 0.0052 - accu
racy: 0.9999 - val_loss: 0.7158 - val_accuracy: 0.8548
Epoch 19/20
30/30 [=====] - 2s 56ms/step - loss: 0.0105 - accu
racy: 0.9977 - val_loss: 0.6848 - val_accuracy: 0.8629
Epoch 20/20
30/30 [=====] - 2s 62ms/step - loss: 0.0031 - accu
racy: 0.9999 - val_loss: 0.7149 - val_accuracy: 0.8649


```

```
Out[11]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [12]: # Training Loss vs Validatin Loss
loss_values = history_dict_onelayer["loss"]
val_loss_values = history_dict_onelayer["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, label="Training loss", color = 'red', linestyle='--', marker='^')
plt.plot(epochs, val_loss_values, label="Validation loss", color = 'green', linestyle='--', marker='^')
plt.title("One Hidden Layer Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
plt.clf()

# Training Accuracy vs. Validation Accuracy
acc = history_dict_onelayer["accuracy"]
val_acc = history_dict_onelayer["val_accuracy"]
plt.plot(epochs, acc, label="Training acc", color='red', linestyle='--', marker='^', mar
plt.plot(epochs, val_acc, label="Validation acc", color='green', linestyle='--', marker='^')
plt.title("One Hidden Layer Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```





Part 2 - Three Layers, 100 Nodes each

```
In [13]: model = keras.Sequential([
    layers.Dense(100, activation="relu"),
    layers.Dense(100, activation="relu"),
    layers.Dense(100, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])

history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))

history_dict_threelayers = history.history
history_dict_threelayers.keys()
```

Epoch 1/20

30/30 [=====] - 4s 94ms/step - loss: 0.5087 - accuracy: 0.7556 - val_loss: 0.3272 - val_accuracy: 0.8741

Epoch 2/20
30/30 [=====] - 2s 78ms/step - loss: 0.2609 - accuracy: 0.8992 - val_loss: 0.3495 - val_accuracy: 0.8590
Epoch 3/20
30/30 [=====] - 2s 60ms/step - loss: 0.1746 - accuracy: 0.9323 - val_loss: 0.2879 - val_accuracy: 0.8842
Epoch 4/20
30/30 [=====] - 2s 51ms/step - loss: 0.1232 - accuracy: 0.9565 - val_loss: 0.3320 - val_accuracy: 0.8834
Epoch 5/20
30/30 [=====] - 2s 54ms/step - loss: 0.0872 - accuracy: 0.9706 - val_loss: 0.4254 - val_accuracy: 0.8792
Epoch 6/20
30/30 [=====] - 2s 54ms/step - loss: 0.0964 - accuracy: 0.9760 - val_loss: 0.4099 - val_accuracy: 0.8834
Epoch 7/20
30/30 [=====] - 2s 51ms/step - loss: 0.0115 - accuracy: 0.9976 - val_loss: 0.5513 - val_accuracy: 0.8809
Epoch 8/20
30/30 [=====] - 2s 54ms/step - loss: 0.1177 - accuracy: 0.9819 - val_loss: 0.5193 - val_accuracy: 0.8824
Epoch 9/20
30/30 [=====] - 2s 52ms/step - loss: 0.0029 - accuracy: 0.9997 - val_loss: 0.6162 - val_accuracy: 0.8814
Epoch 10/20
30/30 [=====] - 2s 61ms/step - loss: 0.0010 - accuracy: 0.9999 - val_loss: 0.7648 - val_accuracy: 0.8810
Epoch 11/20
30/30 [=====] - 1s 50ms/step - loss: 2.4093e-04 - accuracy: 1.0000 - val_loss: 0.9356 - val_accuracy: 0.8799
Epoch 12/20
30/30 [=====] - 1s 50ms/step - loss: 0.0608 - accuracy: 0.9963 - val_loss: 1.6195 - val_accuracy: 0.8239
Epoch 13/20
30/30 [=====] - 2s 55ms/step - loss: 0.0053 - accuracy: 0.9989 - val_loss: 0.9696 - val_accuracy: 0.8801
Epoch 14/20
30/30 [=====] - 2s 51ms/step - loss: 3.0510e-05 - accuracy: 1.0000 - val_loss: 1.0182 - val_accuracy: 0.8792
Epoch 15/20
30/30 [=====] - 2s 55ms/step - loss: 1.9752e-05 - accuracy: 1.0000 - val_loss: 1.0752 - val_accuracy: 0.8789
Epoch 16/20
30/30 [=====] - 1s 50ms/step - loss: 1.1328e-05 - accuracy: 1.0000 - val_loss: 1.1483 - val_accuracy: 0.8791
Epoch 17/20
30/30 [=====] - 2s 58ms/step - loss: 5.1798e-06 - accuracy: 1.0000 - val_loss: 1.2348 - val_accuracy: 0.8779
Epoch 18/20
30/30 [=====] - 2s 60ms/step - loss: 1.8011e-06 - accuracy: 1.0000 - val_loss: 1.3528 - val_accuracy: 0.8780
Epoch 19/20
30/30 [=====] - 2s 50ms/step - loss: 5.0501e-07 -

accuracy: 1.0000 - val_loss: 1.4765 - val_accuracy: 0.8781

Epoch 20/20

30/30 [=====] - 2s 54ms/step - loss: 1.3725e-07 -

accuracy: 1.0000 - val_loss: 1.5845 - val_accuracy: 0.8778

Out[13]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

In [14]: *# Training Loss vs Validatin Loss*

```
loss_values = history_dict_threelayers["loss"]
```

```
val_loss_values = history_dict_threelayers["val_loss"]
```

```
epochs = range(1, len(loss_values) + 1)
```

```
plt.plot(epochs, loss_values, label="Training loss", color = 'red', linestyle='--',
```

```
plt.plot(epochs, val_loss_values, label="Validation loss", color = 'green',
```

```
plt.title("One Hidden Layer Training and validation loss")
```

```
plt.xlabel("Epochs")
```

```
plt.ylabel("Loss")
```

```
plt.legend()
```

```
plt.show()
```

```
plt.clf()
```

Training Accuracy vs. Validation Accuracy

```
acc = history_dict_threelayers["accuracy"]
```

```
val_acc = history_dict_threelayers["val_accuracy"]
```

```
plt.plot(epochs, acc, label="Training acc", color='red', linestyle='--', mar
```

```
plt.plot(epochs, val_acc, label="Validation acc", color='green', linestyle='
```

```
plt.title("One Hidden Layer Training and validation accuracy")
```


```
plt.xlabel("Epochs")
```

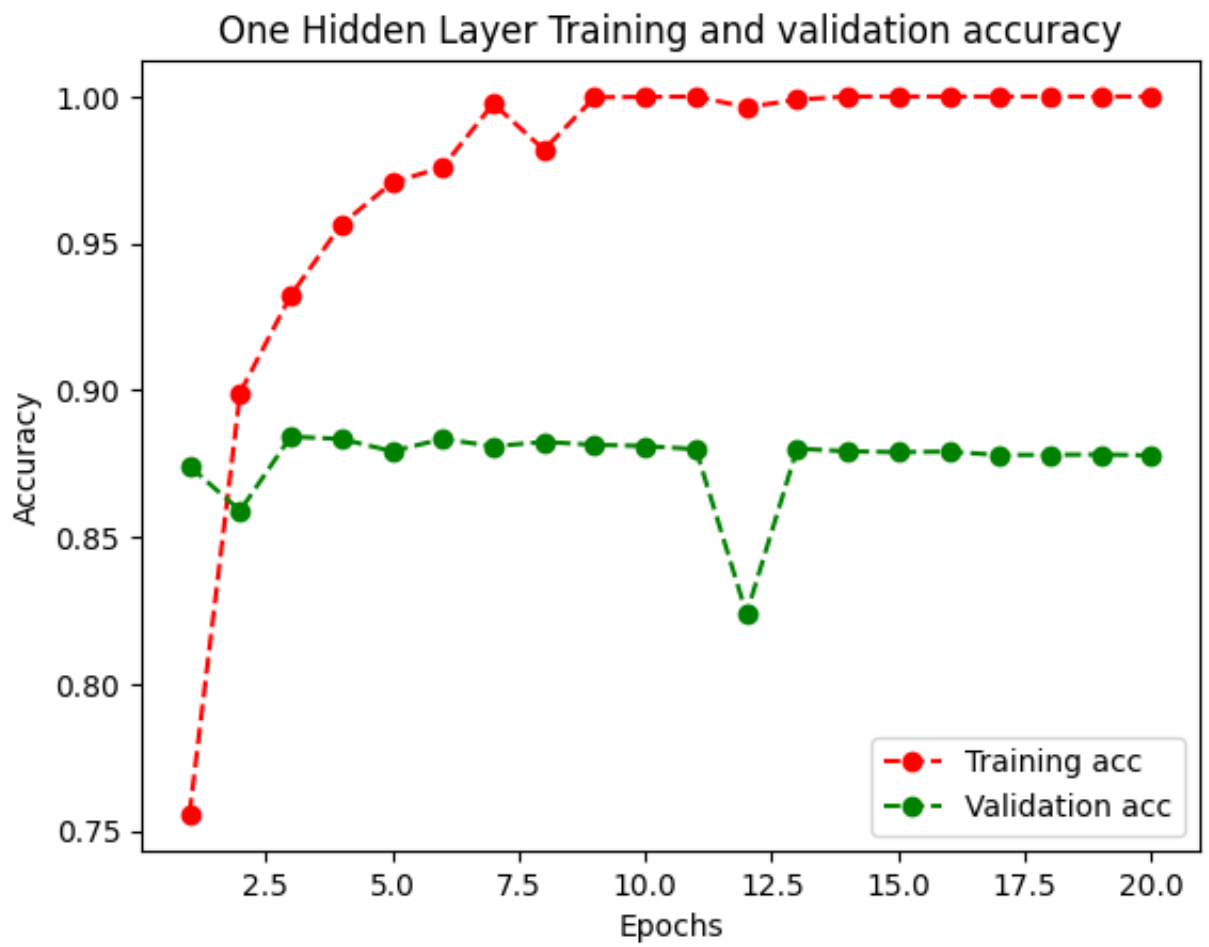
```
plt.ylabel("Accuracy")
```

```
plt.legend()
```

```
plt.show()
```

One Hidden Layer Training and validation loss





Results from Part 1 and Part 2 Plotted Against each other

```

In [15]: loss_values = [
            history_dict_onelayer['loss'],
            history_dict_onelayer['val_loss'],
            history_dict_threelayers['loss'],
            history_dict_threelayers['val_loss']
        ]

acc_values = [
            history_dict_onelayer['accuracy'],
            history_dict_onelayer['val_accuracy'],
            history_dict_threelayers['accuracy'],
            history_dict_threelayers['val_accuracy']
        ]

loss_labels = ['One Layer ', 'Three Layer ', 'Training Loss', 'Validation Loss']
accuracy_labels = ['One Layer ', 'Three Layer ', 'Training Accuracy', 'Validation Accuracy']
colors = ['red', 'green', 'blue', 'magenta']

for idx, x in enumerate(loss_values):
    plt.plot(epochs, x, label=loss_labels[int(idx/2)] + loss_labels[idx%2],
             color=colors[idx], linestyle='--', marker = '^')


plt.title("One Layer vs. Three Layer Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
plt.clf()

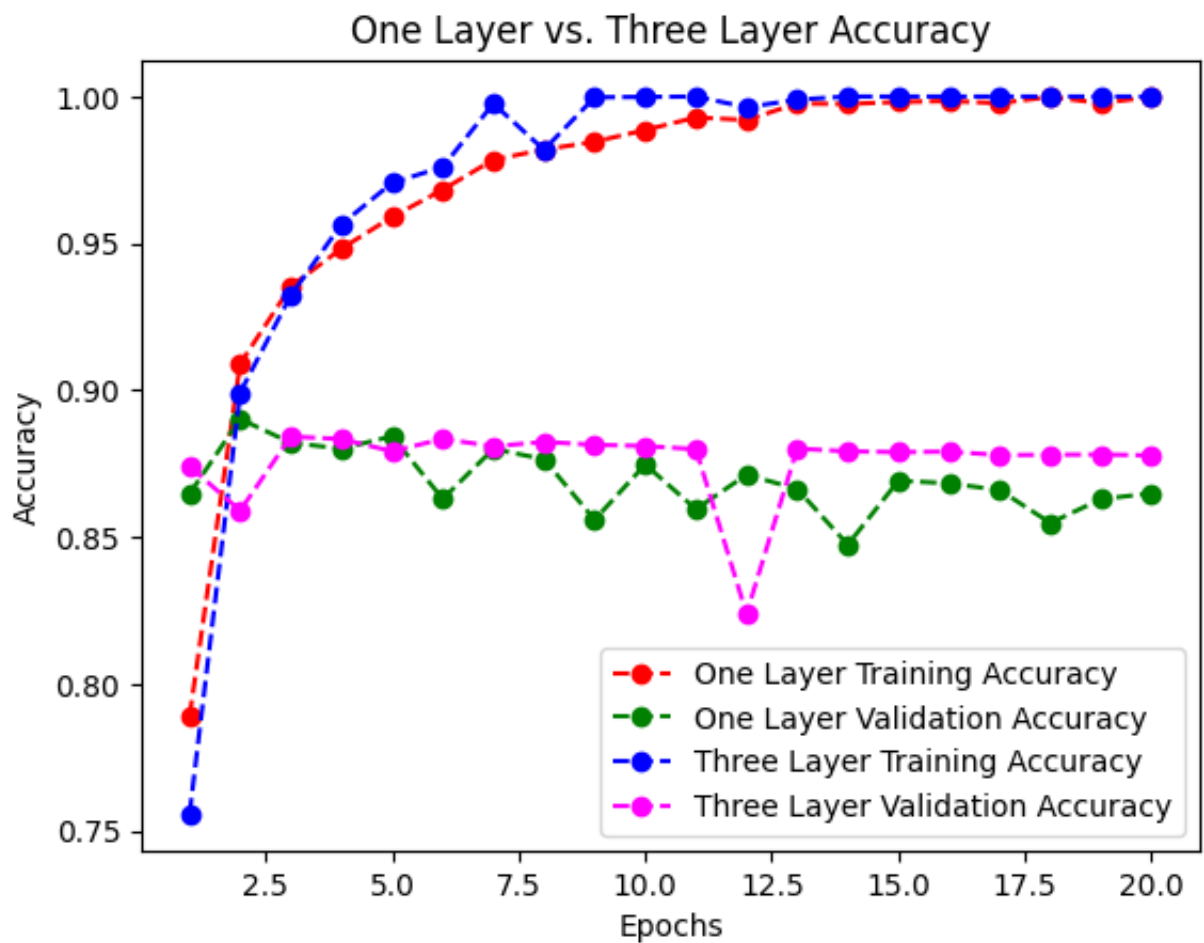
for idx, x in enumerate(acc_values):
    plt.plot(epochs, x, label=accuracy_labels[int(idx/2)] + accuracy_labels[idx%2],
             color=colors[idx], linestyle='--', marker = 'o')

plt.title("One Layer vs. Three Layer Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
plt.clf()

```

One Layer vs. Three Layer Loss





<Figure size 640x480 with 0 Axes>

One of the biggest differences between the 100 node and the 16 node architecture is 100 node architecture plateaus at a much higher validation accuracy while maintaining a similar high training accuracy and corresponding low losses.

Other than that, the differences between both the Three Layer Model and the One Layer model both running the 100 node architecture, the Three Layer Model plateaus at a much higher validation accuracy and correspondingly lower loss.

3) Evaluate the test accuracy of the model before any training has taken place. Compare this to the accuracy post-training, and comment on whether or not this is what you'd expect.

Question attempted with 3 layer model with 100 node architecture as from previous tests, this gave the best results.

Pretraining results

```
In [16]: pretraining_model = keras.Sequential([
          layers.Dense(100, activation="relu"),
          layers.Dense(100, activation="relu"),
          layers.Dense(100, activation="relu"),
          layers.Dense(1, activation="sigmoid")
        ])
pretraining_model.compile(optimizer="rmsprop",
                          loss="binary_crossentropy",
                          metrics=["accuracy"])
pretraining_model.evaluate(x_test, y_test)
```

```
782/782 [=====] - 4s 5ms/step - loss: 0.6938 - accuracy: 0.4966
```

```
Out[16]: [0.6938191652297974, 0.4966000020503998]
```

```
In [17]: posttraining_model = keras.Sequential([
          layers.Dense(100, activation="relu"),
          layers.Dense(100, activation="relu"),
          layers.Dense(100, activation="relu"),
          layers.Dense(1, activation="sigmoid")
        ])
posttraining_model.compile(optimizer="rmsprop",
                          loss="binary_crossentropy",
                          metrics=["accuracy"])
posttraining_model.fit(x_train, y_train, epochs = 4, batch_size=512, validate
```

```
Epoch 1/4
```

```
49/49 [=====] - 14s 261ms/step - loss: 0.4371 - accuracy: 0.8038 - val_loss: 0.3197 - val_accuracy: 0.8698
```

```
Epoch 2/4
```

```
49/49 [=====] - 3s 64ms/step - loss: 0.2379 - accuracy: 0.9075 - val_loss: 0.2920 - val_accuracy: 0.8815
```

```
Epoch 3/4
```

```
49/49 [=====] - 3s 67ms/step - loss: 0.1694 - accuracy: 0.9350 - val_loss: 0.5071 - val_accuracy: 0.8044
```

```
Epoch 4/4
```

```
49/49 [=====] - 3s 65ms/step - loss: 0.1207 - accuracy: 0.9535 - val_loss: 0.3135 - val_accuracy: 0.8796
```


```
Out[17]: <keras.callbacks.History at 0x17a0a89db80>
```


The results kind of represent what I expected. The untrained model achieved a accuracy of 50% which is higher than expected but hopefully the histogram of ouputs provides more insight into this. The trained model achieved a much higher accuracy with an accuracy of 87.5% after 4 epochs of training.

4) Generate a histogram of the output probabilities. Explain what this represents. (use 10 bins)

```
In [18]: pre_prediction = pretraining_model.predict(x_test)
plt.hist(pre_prediction, bins=10)
plt.gca().set(title='Pre Training Histogram');
```


782/782 [=====] - 4s 3ms/step



Post Training Results

```
In [19]: post_prediction = posttraining_model.predict(x_test)
plt.hist(post_prediction, bins=10)
plt.gca().set(title='Post Training Histogram');
```

782/782 [=====] - 2s 3ms/step



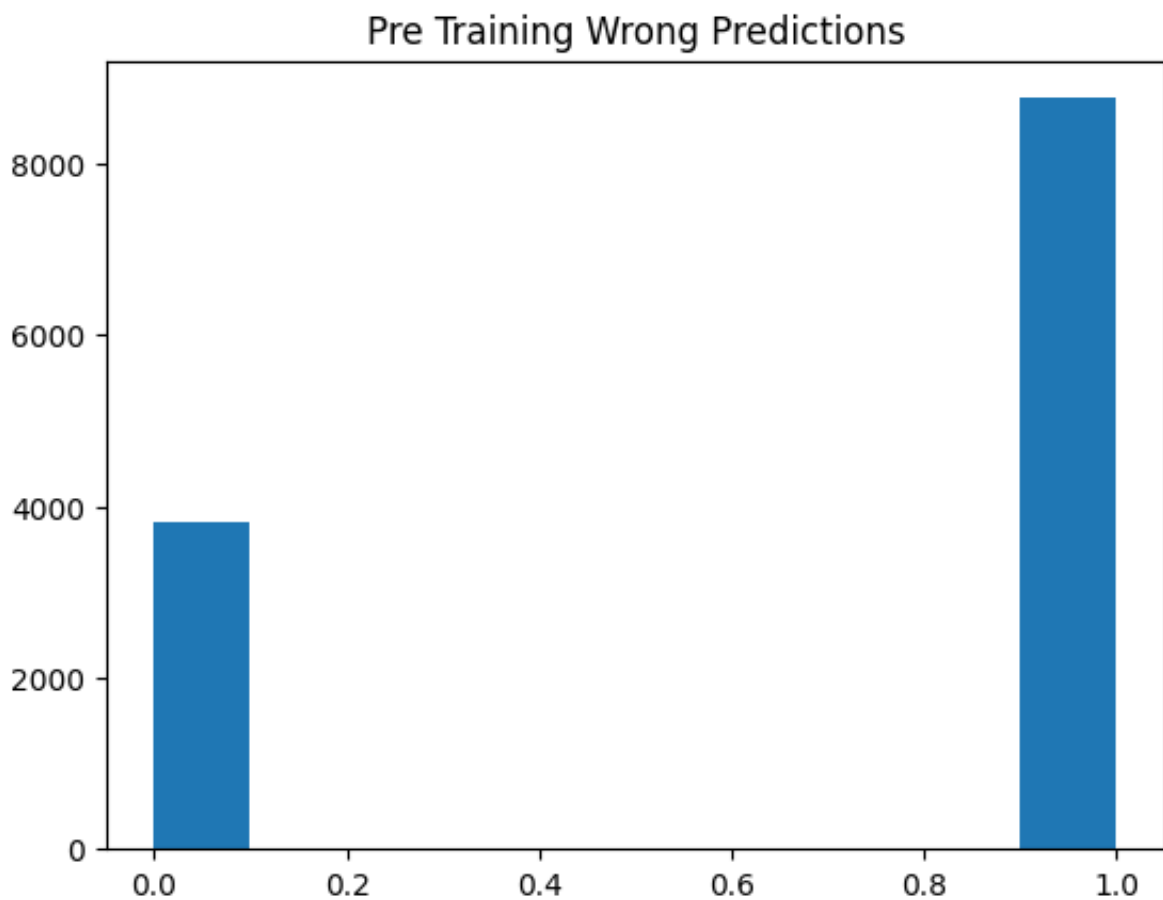
The histogram did provide more insight. In the untrained model, the results are centered roughly around 0.5, in which values below 0.5 were rounded to 0 and values above, to 1. This loosely explains why the results had around a 50% accuracy.

After the model was trained, it provided more concrete predictions as to whether it was a 1 or 0.

5) Find all the predictions where the model got the answer wrong. Generate a histogram of the corresponding output probabilities. Explain what this represents, and comment on whether or not it looks like you'd expect. (Use 10 bins)

```
In [20]: pre_prediction_falses = np.round(pre_prediction)[: , 0]
pre_prediction_falses = pre_prediction_falses[pre_prediction_falses != y_test]
num_untrained_falses = len(pre_prediction_falses)
plt.hist(pre_prediction_falses, bins=10)
plt.gca().set(title='Pre Training Wrong Predictions')
```

```
Out[20]: [Text(0.5, 1.0, 'Pre Training Wrong Predictions')]
```




The pre training model predicted around 12500 results wrong - 12000 of them were predicted to be 1s but were actually 0s and around 500 were 0s when they should have been 1s. This matches our expectations of a 50% accuracy as 12500 is 50% of the original dataset. Moreover we noticed the pre-training histogram was skewed to the right, which means that more values were predicted to be ones and correspondingly a larger number were false predictions.

```
In [23]: post_prediction_falses = np.round(post_prediction)[:, 0]
post_prediction_falses = post_prediction_falses[post_prediction_falses != y_
num_trained_falses = len(post_prediction_falses)
print(len(post_prediction_falses))
plt.hist(post_prediction_falses, bins=10)
plt.gca().set(title='Post Training Wrong Predictions')
```

3010

Out[23]: [Text(0.5, 1.0, 'Post Training Wrong Predictions')]



The Trained Model predicted a total of 3010 values wrong (around 1550 were predicted to be 0s instead of 1s and around 1.5k were predicted to be 1s instead of 0s). 3010 wrong values results in an error rate of 12.04% which corresponds to a accuracy of 87.86% which matches with our results from the sections before. No unexpected behaviour was seen here.

6) Divide the answers from the second histogram by the answers from the first histogram. Explain what this represents.

Dividing the number of errors by the total number of samples ($\text{second_hist}/\text{first_hist}$) gives us the error rate.

```
In [22]: error_rate_untrained = num_untrained_falses/len(pre_prediction)
print("The untrained model's error rate was {}".format(error_rate_untrained))

error_rate_trained = num_trained_falses/len(post_prediction)
print("The trained model's error rate was {}".format(error_rate_trained))
```

The untrained model's error rate was 0.5034
The trained model's error rate was 0.1204

7) Try to improve the model performance. The team with the highest score will receive 50 extra points on the assignment.

```
In [31]: model = keras.Sequential([
    layers.Dropout(0.2),
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dropout(0.2),
    layers.Dense(1, activation="sigmoid")
])

model.compile(optimizer="Adam",
              loss="binary_crossentropy",
              metrics=["accuracy"])

history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=10,
                    batch_size=2000,
                    validation_data=(x_val, y_val))
```

```
Epoch 1/10
8/8 [=====] - 8s 454ms/step - loss: 0.6831 - accuracy: 0.5941 - val_loss: 0.6478 - val_accuracy: 0.7404
Epoch 2/10
8/8 [=====] - 3s 367ms/step - loss: 0.6055 - accuracy: 0.7486 - val_loss: 0.5460 - val_accuracy: 0.8227
Epoch 3/10
8/8 [=====] - 3s 330ms/step - loss: 0.5049 - accuracy: 0.8067 - val_loss: 0.4502 - val_accuracy: 0.8553
Epoch 4/10
8/8 [=====] - 3s 328ms/step - loss: 0.4150 - accuracy: 0.8509 - val_loss: 0.3820 - val_accuracy: 0.8699
Epoch 5/10
8/8 [=====] - 3s 338ms/step - loss: 0.3491 - accuracy: 0.8743 - val_loss: 0.3349 - val_accuracy: 0.8786
Epoch 6/10
8/8 [=====] - 3s 321ms/step - loss: 0.3021 - accuracy: 0.8924 - val_loss: 0.3057 - val_accuracy: 0.8846
Epoch 7/10
8/8 [=====] - 3s 328ms/step - loss: 0.2653 - accuracy: 0.9037 - val_loss: 0.2900 - val_accuracy: 0.8899
Epoch 8/10
8/8 [=====] - 3s 339ms/step - loss: 0.2357 - accuracy: 0.9182 - val_loss: 0.2816 - val_accuracy: 0.8908
Epoch 9/10
8/8 [=====] - 3s 339ms/step - loss: 0.2186 - accuracy: 0.9213 - val_loss: 0.2783 - val_accuracy: 0.8912
Epoch 10/10
8/8 [=====] - 3s 313ms/step - loss: 0.2001 - accuracy: 0.9277 - val_loss: 0.2790 - val_accuracy: 0.8895
```

The highest accuracy we achieved was 89.12 using this model. We used dropout layers which utilizes a mask to ignore the inputs from certain nodes during training. This thereby kind of reduces overfitting

In []: